

Those Delicious Texels

or...

*Dynamic Image-Space Per-Pixel
Displacement Mapping with Silhouette
Antialiasing via Parallax Occlusion
Mapping*



Natalya Tatarchuk
3D Application Research Group
ATI Research, Inc.



Overview of the Talk

- > The goal
- > Overview of approaches to simulate surface detail
- > Parallax occlusion mapping
 - > Theory overview
 - > Algorithm details
 - > Performance analysis and optimizations
 - > In the demo
 - > Art considerations
 - > Uses and future work
- > Conclusions



What Exactly Is the Problem?

- > We want to render very detailed surfaces
- > Don't want to pay the price of millions of triangles
 - > Vertex transform cost
 - > Memory
- > Want to render those detailed surfaces correctly
 - > Preserve depth at all angles
 - > Dynamic lighting
 - > Self occlusion resulting in correct shadowing
- > Thirsty graphics cards want more ALU operations (and they can handle them!)
 - > Of course, this is a balancing game - fill versus vertex transform – you judge for yourself what's best



This is Not Your Typical Parallax Mapping!

- > *Parallax Occlusion Mapping* is a *new* technique
- > Different from
 - > Parallax Mapping
 - > Relief Texture Mapping
- > Per-pixel ray tracing at its core
 - > There are some very recent similar techniques
- > Correctly handles complicated viewing phenomena and surface details
 - > Displays motion parallax
 - > Can render complex geometric surfaces such as displaced text / sharp objects
 - > Surfaces with self-occlusion result in correct self-shadowing
 - > Uses flexible lighting model



We Want the Details... but Not All the Work!... (the early days)

- > First there was bump mapping... [Blinn78]
 - > Detailed and uneven surfaces in some pre-determined manner
 - > Perturbs surface normal using a texture
 - > Popularized as normal mapping – *per-pixel*
- > Doesn't take into account geometric depth of surface
 - > Does not exhibit **parallax**
 - > No self-shadowing of the surface
 - > Smooth silhouettes tell the truth about the surface...





apparent displacement of the object due to viewpoint change



Bump Mapping and Tangent Space

- > Normal vectors are computed in the tangent space for each vertex
- > Tangent space orients the normal map at each vertex
 - > Technically \mathbf{u} and \mathbf{v} can be arbitrarily oriented
 - > Tangent space basis vectors \mathbf{T} and \mathbf{B} orient \mathbf{u} and \mathbf{v}





Lighting in Tangent Space

> To compute lighting from normals in tangent space, we must transform the light vector into tangent space

> *This must be done for each vertex*

> Tangent space is computed at preprocessing time


$$\begin{pmatrix} t_x & t_y & t_z & 0 \\ b_x & b_y & b_z & 0 \\ n_x & n_y & n_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$



Modeling Self-Shadowing


> Interactive Horizon Mapping

- > Based on horizon mapping [Max88]
- > Encodes the height of the shadowing horizon at each point on the bump map in a series of textures for 8 directions
- > Determines the amount of self-shadowing for a given light position
- > At each frame project the light vector onto local tangent plane and compute per-pixel lighting
- > Draw backs:
 - > Texture memory
 - > Original approach: Multipass




[Sloan00]


North Horizon Map



Normal map




North Basis Texture





More Self-Shadowing Efforts

- > Illuminating micro geometry
 - > Precomputes visibility
 - > Computes light scattering in height fields
 - > Uses the Method of Dependent Tests (Monte Carlo technique) to map it to multitexturing hardware



[Heidrich00]



The Holy Grail: Displacement Mapping


- > Tessellates the original surface into a large number of small triangles
 - > Vertices displaced by a height field
 - > RenderMan-quality displacement mapping – sub-pixel triangles
- > Looks great!
 - > Provides self-shadowing / self-occlusion / correct silhouettes
- > Comes at a cost:
 - > Large number of triangles → memory / transform
 - > Not currently available on consumer hardware



View-Dependent Displacement Mapping

- > Per-pixel technique using precomputed surface description (from a given height map)
- > Handles self-shadowing and self-occlusion and correct display of silhouette detail
- > View-dependent approach gives convincing parallax effect
 - > Stores texel relationship from several viewing directions
- > Pros:
 - > Convincing results at high frame rates
 - > Good parallax effects

[Wang03]






View-Dependent Displacement Mapping Costs

- > Multiple texture maps to store per-pixel surface description
- > Preprocessing time: computing the description texture maps from the height map
- > Uses pixel shader: 40 instructions and 14 texture lookups
- > 128 x 128 height fields w/ 32 x 8 view directions = 64Mb uncompressed / 4Mb compressed textures
 - > But 128 x 128 resolution for height / normal map is not typical for complicated objects! It's usually higher



View-Dependent Tessellation of Displacement Maps

[Doggett00]



- > Recursive tessellation on the fly
 - > Uses displacement map to decide tessellation level
- > View-dependent resolution of the displaced mesh
 - > Uses screen size of the current edge to stop tessellation
- > Avoids generating cracks on the displaced surface
 - > Tessellation based only on edge information – not on neighbor information → friendly to current hardware
- > Uses vertex transform units to calculate view-dependence test during tessellation
- > In the end... render many, many triangles



Getting That Feeling... of Motion, of Course

- > Wanted: Motion Parallax
 - > Surface should appear to move correctly with respect to the viewer
- > Recently many approaches appeared to solve this for rendering
 - > Relief Texture Mapping by Oliveira et al in 2000
 - > Parallax Mapping was introduced by Kaneko in 2001
 - > Popularized by Welsh in 2003 with offset limiting technique
 - > Yerex discussed displacement mapping with ray casting in hardware in a Siggraph 2004 sketch



Relief Texture Mapping

- > Supports surface details in 3D and view motion parallax
- > Uses a relief texture – orthogonal displacements per texel
 - > Rendered as a regular texture when viewed from far away
 - > When viewer is close, the mapping is modified (“warped”)
 - > Can be rendered as a mesh of micro polygons when extremely close
- > Uses 1D image operations to apply pre-warp transform to map the relief texture to the polygon
 - > Pre-warp happens in software
 - > Essentially determines the portion of the surface that would be visible when viewed from that angle instead of what is actually visible)
- > Reconstructs the pre-warped texture with two passes to be viewed as a texture mapped polygon



Relief Texture Mapping (cont.)

> Pros:

- > Quality results: correctly handles self-occlusions
- > Produces accurate silhouettes

> Cons:

- > Software operations for image pre-warp
- > Memory footprint for relief textures




[Oliveira00]

- > Additional preprocessing time to compute the relief textures



Parallax Mapping

- > Introduced by Kaneko in 2001
- > Simple way to approximate motion parallax effects on a given polygon
- > Dynamically distorts the texture coordinate to approximate motion parallax effect
- > Shifts texture coordinate using the view vector and the current height map value





Parallax Mapping (cont.)

>Issues:

- > Doesn't accurately represent surface depth
- > Swimming artifacts at grazing angles
- > Flattens geometry at grazing angles

>Pros:

- > No additional texture memory and very quick (~3 extra instructions)



Parallax Mapping with Offset Limiting

- > Same idea as in [Kaneko01], implemented as OpenGL ARB_*_program
- > Uses height map to determine texture coordinate offset for approximating parallax
- > Uses view vector in tangent space to determine how to offset the texels
- > Reduces visual artifacts at grazing angles ("swimming" texels) by limiting the offset to be at most equal to current height value
 - > Flattens geometry significantly at grazing angles
 - > Just a heuristic





Demo

Parallax Mapping with Offset Limiting



Parallax Occlusion Mapping

- > First ideas were introduced in ShaderX³ in 2004 (“Self-Shadowing, Perspective-Correct Bump Mapping Using Reverse Height Map Tracing”, Z. Browley, N. Tatarchuk [Browley04])
- > Current algorithm has several significant improvements over the earlier technique
- > Designed to simulate the effects of per-pixel displacement mapping with correct motion parallax



Demo

Parallax Occlusion Mapping circa 2004



That's All Great, But...

- > We've got even better stuff!!..
 - > New algorithm developed by N. Tatarchuk (aka: me)
- > And that's what we're going to focus on next:



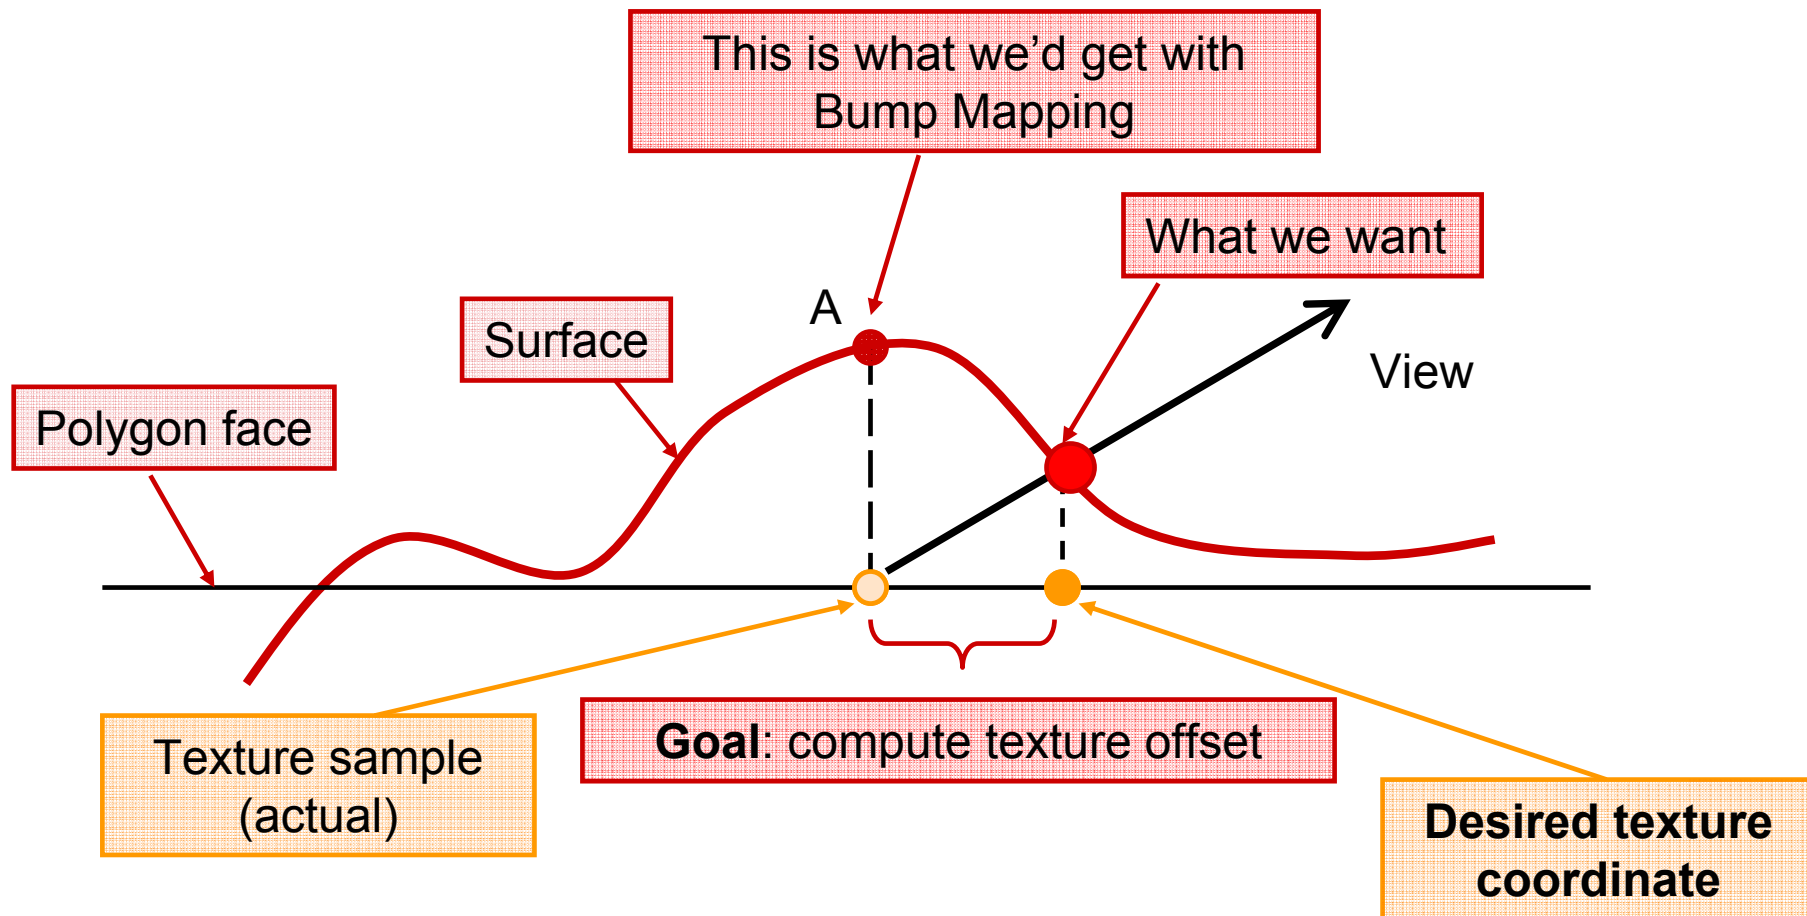


POM Algorithm Overview

- > Compute parallax offset vector
- > Determine the intersection of the height map isosurface with the view ray in tangent space using the reverse parallax offset vector
 - > This is called *Reverse Height Map Tracing*
 - > Computes texture offset coordinates
- > Determine visibility using the light vector in tangent space and currently visible point from above
- > Compute lighting




Watch Those Texels Shift





Reverse Height Field Tracing

- > Ray trace along the reverse parallax vector looking for an intersection with the height field
- > Results in a texture offset corresponding to point of surface visible to the viewer adjusted for parallax





Those Darn Oblique Angles...

- > Any sampling-based algorithm is prone to aliasing
- > We control the sampling rate based on length of the parallax offset vector and surface / view angle relationship
 - > View-dependent sampling rate for correct intersections at grazing angles
 - > Do not sample past the length of the parallax vector
 - > Use $N \bullet V$ to determine the optimal number of samples
 - > This can be artist-controlled by setting min / max samples
 - > Just a simple lerp instruction




Or else!..

- > Otherwise many missed intersections
 - > Previous work [circa 2004] flattened bumps due to perspective biasing
 - > Without perspective biasing, sampling planes were obvious
 - > With view-dependent sampling, no flattening and shape is well-preserved




To Lerp or Not to Lerp: The Comparison: Perspective Bias






To Lerp or Not to Lerp: The Comparison: No Perspective Bias





To Lerp or Not to Lerp: The Comparison: The Right Way





How Does One Render Height Maps, Exactly?

- > Two possibilities
 - > Render surface details as if “pushed down” – the actual polygonal surface will be above the rendered surface
 - > In this case the top (polygon face) is at height = 1, and the deepest value is at 0
 - > Actually push surface details upward (ala displacement mapping)
- > This affects both the art pipeline and the actual algorithm
- > In the presented algorithm, we render the surface pushed down



Is It Perfect?

- > Tradeoffs between speed and quality
 - > Less samples means more possibility for missed features and incorrect intersections
 - > This can result in stair stepping artifacts at oblique angles
 - > Oblique angles pose problems for all similar approaches
- > This approach does not yield correct silhouettes
 - > Not the end of the world: Author art to hide the corner case
 - > Not exactly displacement mapping
 - > Could encode height field description for each texel and use this information to compute silhouettes
 - > Prohibitive cost for games (Might as well use geometry in many cases)



Some Misbehaved Texels

- > Texture-space silhouette edges alias
 - > Especially bad when viewed from further away and animated camera paths
 - > Need to design the algorithm to correct this problem
- > Hardware guys spend time doing better multisampling algorithms.... Is it all for nothing?
 - > Doesn't quite apply here – we are not rendering polygon edges!
 - > The aliasing occurs in texture space – the edges are all in your mind!



Correcting Aliasing Artifacts

- > The texture-space edges occur because we are not performing ray-height map intersections at sub-pixel precision
 - > The problem is very obvious in production quality demos
- > Need to correct this for good results
 - > Determine edges
 - > Jitter rays and compute separate ray-height field intersections
 - > Not enough to just combine the offsets – must compute full new intersections and lighting model for all offsets for good results



Determining Silhouette Edges

- > Want to do this during the run of the algorithm and quickly
 - > Not with post-processing techniques
- > Use height gradient when found a ray-height field intersection
 - > When previous step is a plateau and the next step exhibits large gradient the pixel is at the edge
 - > Very good results
 - > Also can use $N \bullet V$ to improve fidelity
 - > At higher cost





Use The Power of 3.0 Shaders

- > We are performing additional antialiasing work for edge pixels only
- > In a typical scene that means that only at most 10% of your pixels are going through heavier path
- > Only possible with dynamic flow control
 - > Otherwise you always pay the cost



Self-Occlusion and Shadows

- > Apply the same concept as for reverse height map tracing
 - > Use the light vector (in tangent space) to ray-trace from the newly found intersection offset to determine its visibility
- > This works for self-occlusions of the surface
- > Shadows are computed dynamically
 - > Object correctly displays self shadowing effects
- > This must be done for the POM-displaced point (new texture offset)
 - > In order to compute correct visibility results




- [illegible]



If the Start isn't Right?


- > Important to use the parallaxed offset from height field computed earlier
 - > Otherwise shadows will be floating in space





Soft Shadows

- > Can create soft shadows by filtering the visibility samples as you compute them
- > Reduces aliasing artifacts noticeable in many other "fake" displacement mapping approaches
- > Creating soft shadows only adds a few additional ALU ops if you are already computing light visibility for current point
 - > But looks significantly better!





Lighting the Displaced Pixels

- > Use the computed texture coordinate offset to sample desired maps (albedo, normal, detail, etc.)
- > Given those parameters and the visibility information, we can apply any lighting model as desired
 - > Phong
 - > Compute reflection / refraction
 - > Very flexible
- > For many effects, simply diffuse lighting with base texture looks great
 - > Diffuse only suffices for many effects
 - > Glossy specular easily computed – can use gloss maps to reduce specular in the valleys



POM: Great for Scalability

- > Easy LOD scheme for textured surfaces
- > The idea:
 - > Compute the current mip map level in your shader
 - > If rendering from far away, render the low res geometry with bump mapping (Based on LOD level)
 - > Scale the number of samples during height field ray tracing as a function of the current mip map level
 - > As you get closer to the surface, increase the number of samples



Additional Optimizations

- > Use the length of the parallax offset vector to control the sampling rate for ray-tracing
 - > Never sample past what will be visible
- > View-dependent lerping aids depth fidelity
 - > More samples at grazing angles – watch out
 - > May want to restrict viewing angles if necessary



If You Use Post-Processing...

- > May be cheaper to blur edges at post-process time than jitter rays and re-compute full lighting model
 - > Render displaced texels to one buffer and output the edge map as a second buffer
 - > During post-processing, blur the samples from the back buffer for edge pixels



You Can Render Dynamic Height Fields!

- > Great advantage of this algorithm lies in ability to work with dynamically rendered height fields
 - > Render to height fields texture dynamically
 - > Water waves / procedurally generated objects / noise
 - > Apply parallax occlusion mapping computations as regular computations
 - > Compute the normals directly in the shader as necessary
- > Approaches that use precomputed textures (ala distance maps, surface descriptions, etc) cannot support dynamic height field rendering in real-time
- > Combine this algorithm with the computation of wave heights for interaction of parallax occlusion mapped objects with liquids



Does The World Have to Be Flat?

- > Not at all! Columbus said so..
- > Since the computation is in tangent space, we don't have the planarity restriction that many current approaches have
 - > Works equally well on curved objects
 - > Beware of silhouettes
 - > See art considerations later on






What About Problematic Cases?

- > Can handle traditionally hard displacement mapping cases – no restriction for low pass height fields

- > Raised text



- > Fine features





Not Just Low-Detail Height Maps!

- > This means that the height field does not have to be low frequency (aka: low detail)
 - > However, typically for high frequency (high details) fields the ray intersection routine will require higher sampling rate → higher cost
 - > But still compute shadows and self-occlusion tests



The Juicy Details

- > Really takes advantage of the great architecture of current and next-gen GPUs
 - > Balances texture fetches and control flow with ALU load
 - > Flow control
 - > Uses dynamic flow control when supported
 - > Flow control cost is offset by the ALU / texture fetches
 - > ATI Shader Compiler makes aggressive optimizations
- > Easily supports a range of Dx9 hardware targets
 - > Multipass w/ ps_2_0
 - > Single pass in ps_2_b
 - > Single pass dynamic flow control in ps_3_0



PS_2_0 Shader Details

- > Uses static flow control to compute intersections
 - > Compute parallax offset in first pass, output to render target
 - > In second pass computing lighting and shadow term
- > 8 samples in 64 instructions
 - > Performs really well
 - > Doesn't use dynamic number of iterations so the number of samples for height field tracing is constant
 - > Can use more than one pass to sample height map at higher frequencies
 - > 2-3 passes 8 samples each gives very good results
 - > Makes oblique angles look better! 😊



PS_2_b Shader Details

- > Single pass to compute the parallaxed offset, lighting and self-shadowing
- > Uses static number of iterations to compute height field intersections
- > Great performance
- > Use as many samples as needed for your art / scene
 - > Pay in form of instructions



Where is the Goodness...

- > It's in **SM 3.0!**
- > Uses dynamic flow control and early out during ray-tracing operations
 - > Dynamic flow control is tricky – it helps to know what you're doing
 - > Quick path to disassembly is your friend
- > All of the important optimizations / quality improvements happen here
- > Nicely balances ALU ops with control flow instructions and texture fetches
- > Driver Shader Compiler optimizations in action:
 - > A *200 ALU ops* and *32 texture ops* of the disassembled HLSL shader becomes **96 ALU** and **20 texture fetches**
 - > That's 50% quicker!



How Is This Possible?

- > But you thought it wasn't possible to do this at interactive rates
 - > IT IS! *Good hardware architecture helps it*
- > **Cheaper** than rendering same as geometry (remember - less polys)
 - > Naturally, case by case basis
 - > But – to achieve the same fidelity for some of the harder cases (text, fine features) we would need to have mesh resolution to match per-pixel displacement... COSTLY!
 - > Shadow pass for geometry is expensive
 - > Soft shadows – even more
- > Ultimately – content and engine performance profile determine whether to use



Authoring Art for POM

- > Easiest – less detailed height maps with wide features
 - > If rendering bricks or cobble stones, it helps to have wider grout (“valley”) regions
 - > Soft, blurry height maps perform better
- > This algorithm gives the artist control over the range for displacing pixels
 - > This represents the range of the height field
 - > Easily modifiable to get the right look
- > Remember – the algorithm is pushing down, not up
 - > Use this when placing geometry – may need to play the actual geometry higher than planning to render
 - > Height map: white is the top, black is the bottom



When Using Non-contiguous UVs

- > Some complicated objects may require the use of non-contiguous UVs
- > If not careful, this may result in a visible seam at run time






Removing the Seam

- > This is an easily solvable problem – when authoring texture maps, add padding of ~10-20 pixels around the border for intersecting regions
 - > Important to do this process for *all* texture assets used to render this object (Normal maps / albedo / detail, etc)




Padding the Maps: Original Art



This region caused the seam




Padding the Maps: Intermediate Step





Padding the Maps: Final Art



Notice the
difference



What About Those Silhouettes?

- > Can hide smooth silhouettes with special silhouette geometry
 - > A row of bricks on the corners of the building
 - > Curb of the sidewalk





A Few More Hints

- > For very detailed, complex surfaces, height map / normal map textures need to have enough resolution
- > If computing specular term use gloss maps to reduce specularity in the valleys



Demo

Parallax Occlusion Mapping



Uses and Future Ideas

- > Everybody gets bumpy tonight!
 - > Bumpy surfaces galore - Bricks, cobble stones, manhole covers, grates, you name it!
- > Water waves animation
- > Water moving over occlusion-mapped objects
- > Bullet holes animation
 - > Animating POMapped objects is not a problem, in fact works very well
- > Many many more...



Conclusions

- > *Parallax Occlusion Mapping* is a powerful technique for rendering complex surface details in real time
 - > Image-space algorithm
- > Has modest texture memory footprint
 - > Normal map and height map (tile-able)
- > Produces excellent lighting results
 - > Self-shadowing for self-occlusion in real-time
 - > Soft dynamic shadows
 - > Flexible lighting model
 - > Antialiased texture-space edges
- > Great LOD rendering technique for textured scenes
- > Easily supports dynamic rendering of height fields
- > Efficiently uses existing pixel pipelines for highly interactive rendering
 - > Well-designed for future hardware architectures



References

- > [Blinn78] Blinn, James F. "*Simulation of wrinkled surfaces*", Siggraph '78
- > [Max88] N. Max "*Horizon Mapping: shadows for bump-mapped surfaces*", The Visual Computer 1988
- > [Sloan00] P-P. Sloan, M. Cohen, "*Interactive Horizon Mapping*", Eurographics 2000
- > [Welsh04] T. Welsh, "*Parallax Mapping with Offset Limiting: A Per Pixel Approximation of Uneven Surfaces*", 2004
- > [Wang03] L. Wang *et al*, "View-Dependent Displacement Mapping", Siggraph 2003
- > [Doggett00] M. Doggett, J. Hirche, "*Adaptive View Dependent Tessellation of Displacement Maps*", Eurographics Hardware Workshop 2000
- > [Kaneko01] Kaneko *et al.*, "Detailed Shape Representation with Parallax Mapping", ICAT 2001



More References

- > [Brawley04], Z. Brawley, N. Tatarchuk, "*Parallax Occlusion Mapping: Self-Shadowing, Perspective-Correct Bump Mapping Using Reverse Height Map Tracing*", ShaderX³, 2004
- > [Oliveira00] M. Oliveira et al, "*Relief Texture Mapping*", Siggraph 2000
- > [Yerex04] K. Yerex, M. Jagersand, "*Displacement Mapping with Ray-casting in Hardware*", Siggraph 2004 Sketches
- > And some very latest references
 - > W. Donnelly, "*Per-Pixel Displacement Mapping with Distance Functions*", GPU Gems2, 2005



Acknowledgements

- > Zoe Brawley, Relic Entertainment
- > 3D Application Research Group